

# DevOps Orchestra: A Multi-Agent AI-Powered DevOps Automation System

## Purpose:

Automate the **entire DevOps lifecycle** using **AI agents** that collaborate like a team to handle code validation, infrastructure provisioning, testing, deployment, monitoring, and recovery—without requiring manual scripts or human intervention.

## How It Works

1. **User pushes code to GitHub** along with a `devops_orchestra.yaml` file that defines deployment preferences (e.g., cloud, environment, strategy).
2. The system's **agents** kick in:
  - **GitOps Agent** detects the push via webhook and triggers the workflow.
  - **Code Analysis Agent** inspects code quality and completeness.
  - **Testing Agent** auto-generates and runs unit tests, integration tests using AI.
  - **Build Agent** checks for `Dockerfile` (generates if missing). It builds the code and pushes the Docker image to Docker Hub/ Amazon ECR.
  - **IaC Agent** provisions cloud infra using Terraform (generates if missing).
  - **Deployment Agent** deploys using strategies like blue-green or canary.
  - **Observability Agent** monitors app health and auto-rolls back if needed.
  - **ChatOps Agent** communicates with users via Slack.
3. Agents communicate and reassign tasks among themselves via the shared event bus (Kafka).

## What Makes It Special

- Uses **open-source LLMs** (e.g., LLaMA2, Mistral) to generate missing DevOps artifacts.
- Works even when the user only provides **partial code** (no Docker/Terraform).
- Smart enough to **auto-heal**, **self-correct**, or **ask for user intervention** when needed.
- Supports **dynamic coordination** between agents for efficiency.
- Ideal for a **cloud-native** app.

## Use-Cases:

### 1. Push-to-Deploy: Full Automation on Code Push

**Actors:** Developer, GitOps Agent, Code Analysis Agent, Testing Agent, Build Agent, IaC Agent, Deployment Agent, Observability Agent.

**Scenario:** A developer pushes new backend code to GitHub. DevOps Orchestra autonomously builds, tests, provisions, deploys, and monitors the application.

**Flow:**

1. GitOps Agent receives the GitHub webhook on code push.
  - Parses the devops\_orchestra.yaml file (if present) or uses the default configuration.
  - Publishes a code\_push event to the event bus.
2. Code Analysis Agent listens for the code\_push event.
  - Lints and analyzes the code for errors or missing modules.
  - Publishes code\_analysis\_result event (with pass/fail + warnings).
3. Testing Agent listens for code\_analysis\_result.
  - If valid, generates and runs unit/integration tests.
  - Publishes test\_results event (pass/fail, coverage report).
4. Build Agent listens for test\_results.
  - If tests passed, builds Docker image (or generates Dockerfile if missing) and pushes it to Docker Hub/ AWS ECR.
  - Publishes build\_ready event with image URL or artifact.
5. IaC Agent listens for build\_ready.
  - Checks for existing Terraform/IaC files.
  - Generates and applies infrastructure code (EC2, RDS, etc.).
  - Publishes iac\_ready event with provisioned resources info.
6. Deployment Agent listens for iac\_ready.
  - Deploys the new version using blue-green strategy (or as specified).
  - Publishes deployment\_triggered or deployment\_failed.
7. Observability Agent listens for deployment\_triggered.
  - Monitors health, latency, error rates.
  - If failure is detected, triggers auto rollback.
  - Publishes observability\_alert or rollback\_event.
8. ChatOps Agent listens to all events.
  - Notifies the developer in Slack about:
    - Linting/test results
    - Build status
    - Infra provisioning
    - Deployment success/failure
    - Rollback or alerts

## 2. Code Quality Gate Before Build

**Actors:** Developer, Code Analysis Agent, ChatOps Agent

**Scenario:** Code is pushed but contains anti-patterns or incomplete modules.

**Flow:**

- GitOps Agent triggers the workflow.
- Code Analysis Agent reviews syntax, structure, and style.
- Flags issues like missing modules or inconsistent structure.
- ChatOps Agent sends feedback via Slack with suggestions.
- Deployment halts until issues are fixed.

## 3. Missing Dockerfile – Build Agent Intervention

**Actors:** Build Agent, ChatOps Agent

**Scenario:** The Project lacks a Dockerfile.

**Flow:**

- Build Agent inspects the codebase.
- Based on language/framework (e.g., Node.js, Django), generates a Dockerfile using AI.
- Sends it to the global state (Kafka) for other agents to consume.
- Alerts user via ChatOps Agent.

## 4. Missing Terraform Files – IaC Agent Bootstrapping

**Actors:** IaC Agent, ChatOps Agent

**Scenario:** The repository doesn't contain `main.tf` or infrastructure config.

**Flow:**

- IaC Agent generates Terraform scripts from `devops_orchestra.yaml`.
- Provisions cloud infra (e.g., S3, EC2, Lambda, VPC) based on configuration.
- Stores generated files and alerts via ChatOps Agent.

## 5. Intelligent Test Generation for Legacy Code

**Actors:** Code Analysis Agent, Testing Agent, ChatOps Agent

**Scenario:** An old codebase lacks test coverage.

**Flow:**

- Code Analysis Agent maps out the code structure.
- Testing Agent uses AI to generate unit/integration tests.

- Runs tests in an isolated container.
- Sends test reports to Slack via ChatOps Agent.

## 6. Automated Canary Release With Auto Rollback

**Actors:** Deployment Agent, Observability Agent, ChatOps Agent

**Scenario:** A New version of the app is released.

**Flow:**

- Deployment Agent performs canary release (small % of traffic gets new version).
- Observability Agent monitors metrics (latency, error rates).
- If issues detected, rollback is triggered automatically.
- ChatOps Agent alerts team.

## 7. Chat-Based Deployment Command

**Actors:** ChatOps Agent, Deployment Agent

**Scenario:** Dev wants to deploy a hotfix from Slack.

**Flow:**

- Developer types `/deploy frontend hotfix-v2` in Slack.
- ChatOps Agent parses the intent.
- Sends deployment instruction to Deployment Agent.
- Agent fetches latest Git revision, builds, and deploys.
- Reports back in the same Slack thread.

## 8. Real-Time Issue Detection and Healing

**Actors:** Observability Agent, Deployment Agent, ChatOps Agent

**Scenario:** A frontend microservice crashes due to memory leaks.

**Flow:**

- Observability Agent detects anomalous CPU/memory usage.
- Notifies global event bus.
- Deployment Agent rolls back to the last healthy version.
- ChatOps Agent summarizes crash logs and notifies team.

## 9. Pre-Merge Deployment Validation

**Actors:** GitOps Agent, Code Analysis Agent, Testing Agent, ChatOps Agent

**Scenario:** A pull request is raised for staging deployment.

**Flow:**

- GitOps Agent triggers on PR.
- Code Analysis Agent inspects code for risks.
- Testing Agent generates and runs unit/ integration tests.
- If all tests passes, ChatOps Agent posts a PR comment: “Safe to merge and deploy.”

## 10. Frontend and Backend Coordinated Deployments

**Actors:** GitOps Agent, Build Agent, Deployment Agent

**Scenario:** Frontend and backend are in the same repo.

**Flow:**

- GitOps Agent watches for commits in both `/frontend` and `/backend`.
- Build Agent creates two separate images and pushes them to Docker Hub/ Amazon ECR.
- Deployment Agent coordinates rollout such that frontend and backend are deployed together using `depends_on` logic in `devops_orchestra.yaml`.

## **devops\_orchestra.yaml Template:**

# Required: Project metadata

project:

name: my-app

language: python

framework: flask

repo: <https://github.com/user/my-app>

# Required: Build settings

build:

tool: docker # options: docker, maven, npm, poetry, custom

context: .

dockerfile: Dockerfile

# Required: Test configuration

testing:

enabled: true

framework: pytest # options: pytest, unittest, jest, mocha, etc.

command: pytest tests/

# Required: Deployment

deployment:

type: cloud # options: cloud, on-premise, hybrid

provider: aws # options: aws, gcp, azure, custom

strategy: blue-green # options: blue-green, canary, rolling

region: us-west-2

services:

- name: api-server

runtime: ec2 # options: ec2, lambda, kubernetes, etc.

port: 5000

env:

- key: ENV

value: production

- key: DB\_URL

value: postgres://...

# Required: Infrastructure as Code (IaC)

infrastructure:

tool: terraform # options: terraform, cloudformation, ansible

path: infra/ # path to IaC scripts

# Optional: Secrets (recommend using secret manager)

secrets:

manager: aws\_secrets\_manager

keys:

- DB\_PASSWORD
- API\_KEY

# Optional: Rollback policy

rollback:

enabled: true

threshold:

cpu: 90

errors: 5

duration: 5m

# Optional: Observability

observability:

enabled: true

tools:

- prometheus
- grafana
- cloudwatch

alerts:

- type: latency

threshold\_ms: 500

- type: error\_rate

threshold: 2%

# Optional: ChatOps settings

chatops:

enabled: true

platform: slack

channel: "#devops-orchestra"

notify\_on:

- deployment\_success
- deployment\_failure
- rollback\_triggered

Here are the key **assumptions** made in the **DevOps Orchestra** system design:

## System-Level Assumptions

### 1. GitHub as Source of Truth

All project codebases are stored in GitHub repositories. No support for other VCS like GitLab/Bitbucket unless extended.

### 2. Presence of `devops_orchestra.yaml`

Users must include a configuration file in the root of their repository specifying:

- Environment (cloud/on-prem)
- Infrastructure preferences
- Deployment strategy

*If missing, default behavior or agent intervention kicks in.*

### 3. LLMs Available Locally

The system uses open-source LLMs (like LLaMA 2, Code LLaMA, or Mistral) instead of commercial APIs like GPT-4 due to cost constraints.

### 4. Kafka or Event Bus Available

A working event bus (e.g., Apache Kafka or Redis Streams) is assumed to be set up to facilitate inter-agent communication.

### 5. Microservice-Friendly Codebase

The system assumes the code is modular (e.g., frontend, backend) so that agents can independently build and deploy parts of the system.

## Agent-Level Assumptions

### 1. Build Agent Assumptions

- Can correctly identify the language/framework (Node, Django, Flask, etc.).
- Has access to a template library for Dockerfile generation.

### 2. IaC Agent Assumptions

- The YAML contains sufficient information (cloud provider, region, service type) for Terraform file generation.
- Users are okay with basic default infrastructure if configuration is minimal.



### 3. Code Analysis Agent Assumptions

- Code is at least syntactically correct and the project structure is identifiable.
- Errors can be detected using static analysis + LLM heuristics.

### 4. Testing Agent Assumptions

- Can generate tests using AI even without user-provided ones.
- Enough function/method structure is available in the code for test generation to be meaningful.

### 5. Deployment Agent Assumptions

- Code has been containerized or can be containerized before deployment.
- Infra provisioning is already done successfully before it starts.

### 6. Observability Agent Assumptions

- Application logs, metrics, and health endpoints (like `/health`, `/metrics`) are available or inferred.
- Sufficient permissions to perform rollback or restart actions.

### 7. ChatOps Agent Assumptions

- User has integrated Slack via webhook or bot.
- Users can be identified and authenticated within the chat platform.

## User Assumptions

### 1. Basic YAML Proficiency

Users can define the required fields correctly in `devops_orchestra.yaml`. Defaults are used if some fields are missing.

### 2. Push Triggers Deployment

Users are aware that pushing to certain branches (like `main`, `release`) automatically triggers the DevOps pipeline.

### 3. Cloud Credentials Available

If deploying to cloud, credentials (via GitHub Secrets or environment variables) are accessible to the IaC and Deployment Agents.

#### 4. Failure Handling Is Acceptable

The system may skip or rollback if:

- Infrastructure provisioning fails
- Tests fail
- App health degrades after deployment

### LLM Behavior Assumptions

#### 1. Sufficient Context Window

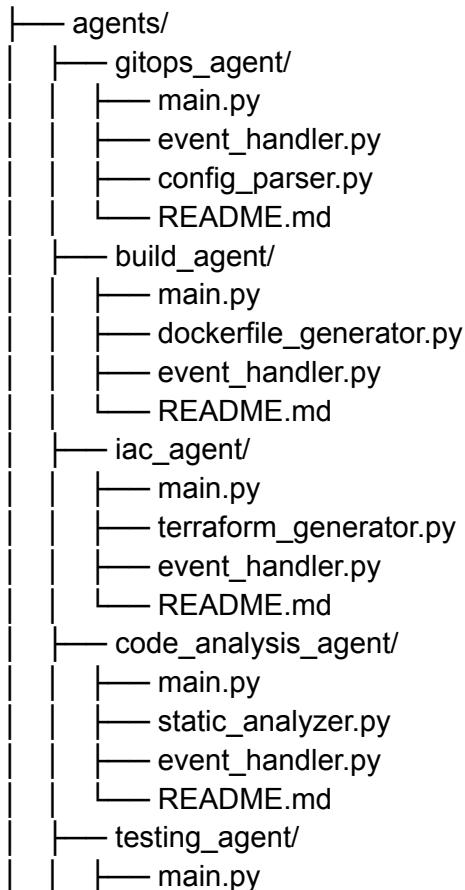
LLMs can process a representative portion of the codebase, but may not handle extremely large monorepos in one go.

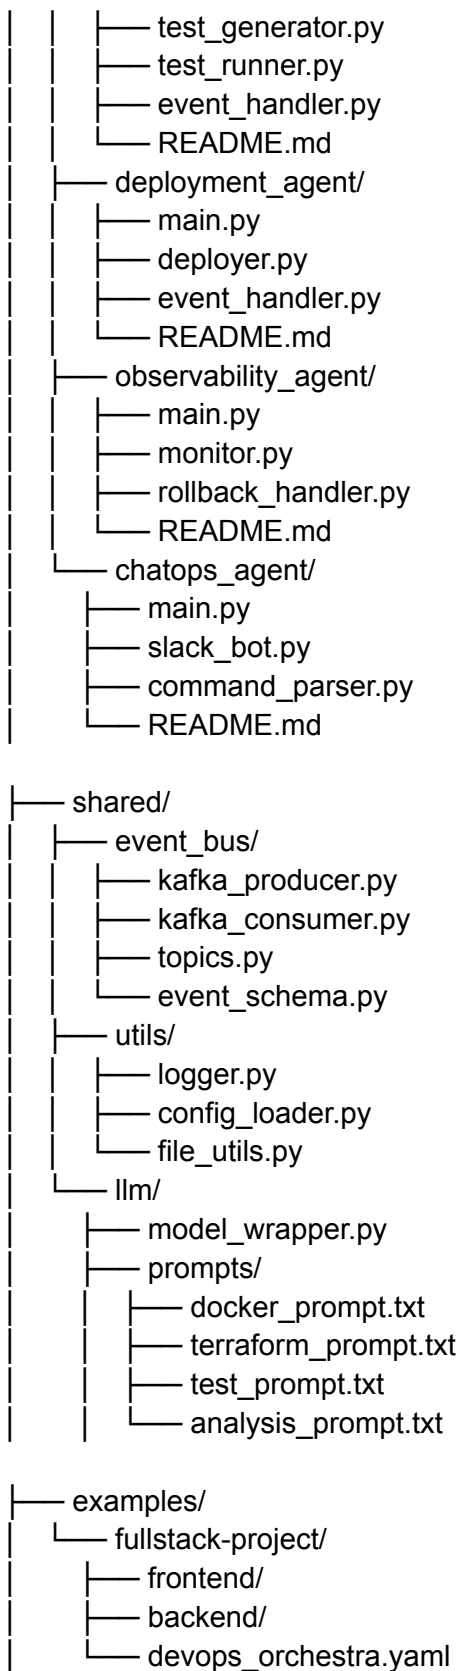
#### 2. Trust in AI-Generated Artifacts

The system assumes the AI-generated Dockerfile, Terraform scripts, and test cases are safe defaults unless overridden.

### Full Project Structure for Peer-to-Peer DevOps Orchestra:

devops-orchestra/





```
|— infrastructure/
|   |— generated/
|   |   |— main.tf
|   |   |— ec2.tf
|   |   |— outputs.tf
|   |— base/
|   |   |— provider.tf
|   |   |— variables.tf
```

```
|— tests/
|   |— test_gitops_agent.py
|   |— test_build_agent.py
|   |— test_testing_agent.py
|   |— test_event_bus.py
|   |— test_utils.py
```

```
|— docker-compose.yml
|— requirements.txt
|— devops_orchestra.yaml # Optional default config
|— README.md
```

---

## Root Directory (**devops-orchestra/**)

### **README.md**

- Project overview and setup guide
- Contains: agent roles, how to run Kafka + agents, usage examples
- **Essential for onboarding and documentation**

---

### **requirements.txt**

Lists common Python dependencies.

---

### **docker-compose.yml**

- Defines and spins up:
  - Kafka

- Zookeeper
  - Agents (e.g., GitOps Agent, Build Agent)
  - Used for **local development and testing**
- 

### **devops\_orchestra.yaml**

- Sample orchestration config (if user repo doesn't provide one)
  - Used to simulate deployments or provide fallback defaults
- 

### **agents/ (Contains One Folder per Agent)**

Example: **gitops\_agent/**

#### **main.py**

- Entry point for the GitOps Agent
- Starts the webhook server or event listener
- Subscribes to push events and publishes **code\_push** events to Kafka

#### **event\_handler.py**

- Handles incoming GitHub push events
- Parses branch, repo, changed files

#### **config\_parser.py**

- Reads and validates the **devops\_orchestra.yaml** file
- Sends its parsed output to the event bus

## README.md

- Explains GitOps Agent's responsibilities and usage
- 

All other agent folders (e.g., `build_agent/`, `iac_agent/`, etc.) follow the same structure:

### `main.py`

- Initializes Kafka consumer for that agent's topic
- Subscribes to relevant events like `code_push`, `build_ready`, `infra_ready`

### `*_generator.py`, `*_handler.py`, etc.

- Agent-specific logic:
    - **Build Agent:** `dockerfile_generator.py` creates Dockerfiles using LLM
    - **IaC Agent:** `terraform_generator.py` creates Terraform files
    - **Testing Agent:** `test_generator.py` creates tests, `test_runner.py` executes them
    - **Deployment Agent:** `deployer.py` performs blue-green/canary deployments
    - **Observability Agent:** `monitor.py` reads logs/metrics, `rollback_handler.py` performs rollback
    - **ChatOps Agent:** `slack_bot.py` connects to Slack, `command_parser.py` maps commands to actions
- 

### `shared/`

### `event_bus/`

### `kafka_producer.py`

- Publishes events to Kafka topics like:
  - `build_ready`
  - `deployment_triggered`

- `rollback_event`

### `kafka_consumer.py`

- Subscribes to one or more Kafka topics
- Dispatches incoming events to appropriate handlers

### `topics.py`

- Defines standard Kafka topic names to ensure consistency across agents

### `event_schema.py`

- Defines and validates event structure (e.g., JSON schema) to ensure clean inter-agent communication

---

### `utils/`

#### `logger.py`

- Standard logging setup for all agents
- Formats logs with timestamp, agent name, and log level

#### `config_loader.py`

- Loads and validates YAML config files (`devops_orchestra.yaml`)

#### `file_utils.py`

- Utilities to check if a file exists, read/write files, copy templates
-

**llm/**

**model\_wrapper.py**

- Abstraction layer for calling LLaMA, Mistral, etc.
- Supports both local inference and Hugging Face API
- Used by agents like Build, IaC, and Testing

**prompts/**

- Stores reusable prompt templates

**Examples:**

- **docker\_prompt.txt:**  
"Generate a Dockerfile for a Flask app with port 5000..."
- **terraform\_prompt.txt:**  
"Generate Terraform code to deploy an EC2 instance..."
- **test\_prompt.txt:**  
"Generate unit tests for this Python function..."
- **analysis\_prompt.txt:**  
"Perform code quality analysis for this React component..."

---

**examples/**

**fullstack-project/**

**frontend/ and backend/**

- Sample app code (React + Flask, etc.)



## devops\_orchestra.yaml

- Real config used to trigger the full DevOps pipeline.
- 

## infrastructure/

### generated/

- Contains Terraform files **generated by the IaC Agent** on the fly

#### Files:

- `main.tf`, `ec2.tf`, `outputs.tf`: auto-generated based on YAML config

### base/

- Contains reusable modules or manually written Terraform templates
- Used as fallback or as scaffolding for generation

#### Files:

- `provider.tf`: AWS provider config
  - `variables.tf`: shared variables for all infra
- 

## tests/

Contains **unit and integration tests** for agents and shared code.

#### Files:

- `test_gitops_agent.py`: Tests for GitOps event handling
- `test_build_agent.py`: Dockerfile generation, edge cases
- `test_testing_agent.py`: AI test generation, runner logic
- `test_event_bus.py`: Kafka message flow
- `test_utils.py`: Utility module test.

Run with: `pytest tests/`